

# Software infrastructure and algorithms to facilitate co-location of observation and model data



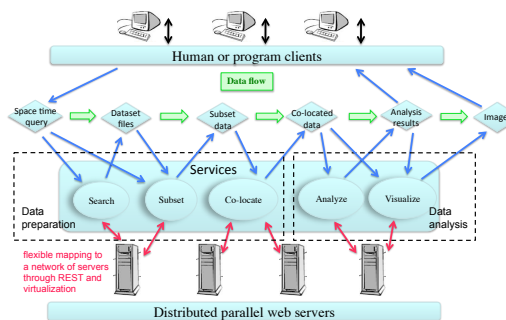
Lei Pan\*, Seungwon Lee, Gary Block, Robert Morris, Jui-lin Li, and Duane Waliser  
Jet Propulsion Laboratory, California Institute of Technology, Pasadena, CA 91101

\*lei.pan@jpl.nasa.gov

## Background and Goal

- Observation data (e.g., from A-Train satellites) and model data (e.g., ECMWF) live in their own respective grid space, have different sampling characteristics, and use different formats and structures in archiving
- Scientists expect synergistic usage of these datasets, e.g., to calibrate one instrument using another, or to conduct model data assimilation using observation
- Co-location, which is to interpolate from a source data grid onto that of a target, provides a bridge for the above gap
- To build a software tool to facilitate co-location
- To make it easy to use and fast by leveraging technologies (e.g., web services, virtualization, and parallelization), and implementing efficient and scalable core algorithms

## Functional Architecture



## Web Services and Virtualization

- Reason for services: huge datasets from satellite and model in PBs
- Open source web services solutions rich and simple: Flask (Python microframework), Gunicorn (Python WSGI server), Tornado (scalable web server)
- Separate app. traffic (Gunicorn) from static HTTP traffic (Tornado)
- RESTful interfaces using ROA: URI to carry "scoping info." and HTTP to convey "method info."; stateless for parallel deployment
- Virtualization using VMWare: service replication made easy



## Parallelization

- Co-location of multiple data granules inherently data parallel
- Python Multiprocessing for multicore:
  - Multicore machines affordable and available
  - Shared memory programming relatively easy
  - Effective memory recycling when child processes join
  - Limited scalability
- Parallel Python for cluster (under construction):
  - Scalable to 100s of cores and beyond
  - Remote data shipping either via parallel file system (e.g., PVFS) or using pre-fetching




## Core Algorithm: Nearest Neighbor Search

Our algorithm has a linear complexity of  $O(n+m)$ , where  $n$  is the number of target points and  $m$  the number of source points


1. Overlay  $k$  by  $k$  grid on target points where  $k = \sqrt{n}$  ( $O(k)$ )
2. Register each target point to a grid cell ( $O(n)$ )
3. For each source point  $S$  ( $O(m)$ )
  1. find minimum bounding rectangle (MBR) of its footprint
  2. Find grid cells that cover MBR
  3. For each grid cell  $G$  ( $O(C)$ )
    - For each target point  $T$  in  $G$ 
      - store  $S$  to  $T$ 's candidate list (sorted by distance= $|T-S|$ )
4. For each target point  $T$  ( $O(n)$ )
  1. Compute average among the first  $L$  in  $T$ 's candidate list

X: target point, O: Source point



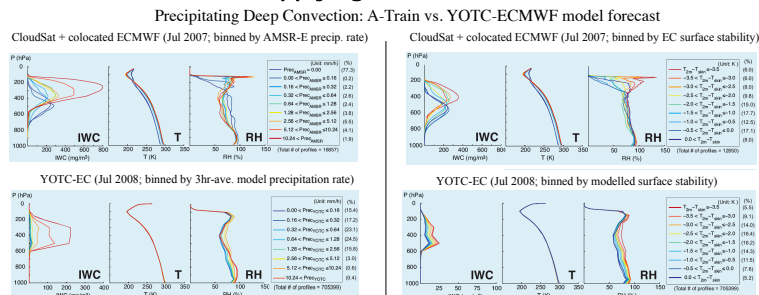
- Nearest neighbor not necessary in same grid cell
- C: a constant bounded by max "density" of target pts in MBR, which is not proportional to  $n$
- When  $L=1$ , nearest neighbor; when  $L>1$ , average of  $L$  nearest neighbors

## Co-location Performance



- nearest neighbor search algorithm improves computational performance by 100-1000 times over the naive  $O(N*N)$  algorithm
- parallel co-location implementation has a speedup of 5 on 8 cores

## Applying to Science



A comprehensive A-Train-ECMWF co-located data set can be used to characterize deep convection cloud structures, properties, and its environmental context. Reasonable agreements between observations and YOTC-EC model forecast are found.

## Conclusions and Future Work

- Web services efficient: large data is where heavy number crunching is, along with powerful CPUs; small result is downloaded by clients; cost of coarse communication overshadowed by gain; light-weight clients both in terms of processing power and number of software packages
- Web services easy to use: similar to function calls that scientists are familiar with, using a browser or a python program, with samples provided
- Efficient algorithm & data structure and parallelization both important for performance
- Demonstrated value of the co-located datasets to scientific research
- ATrain/ECMWF co-located data products released to public in 2011
- Future work: use Parallel Python for scalable distributed computing
- Future work: build data analysis web services (funded by NASA ROSES 2011)

**Acknowledgements:** This work was performed at the Jet Propulsion Laboratory, California Institute of Technology, under contract with NASA.

Presented at AGU Fall Meeting, San Francisco, Dec 3-7, 2012